

Voting Theory in the Lean Theorem Prover

First Author¹[0000–1111–2222–3333], Second Author^{2,3}[1111–2222–3333–4444], and
Third Author³[2222–3333–4444–5555]

¹ Princeton University, Princeton NJ 08544, USA

² Springer Heidelberg, Tiergartenstr. 17, 69121 Heidelberg, Germany
lncs@springer.com

<http://www.springer.com/gp/computer-science/lncs>

³ ABC Institute, Rupert-Karls-University Heidelberg, Heidelberg, Germany
{abc,lncs}@uni-heidelberg.de

Abstract. There is a long tradition of fruitful interaction between logic and social choice theory. In recent years, much of this interaction has focused on computer-aided methods such as SAT solving and interactive theorem proving. In this paper, we report on the development of a framework for formalizing voting theory in the Lean theorem prover, which we have applied to verify all of the claimed properties of a recently proposed voting method. While previous applications of interactive theorem proving to social choice (using Isabelle/HOL and Mizar) have focused on the verification of impossibility theorems, we aim to cover a variety of results ranging from impossibility theorems to the verification of properties of specific voting methods (e.g., Condorcet consistency, independence of clones, etc.). In order to formalize voting theoretic axioms concerning adding or removing candidates and voters, we work in a variable-election setting whose formalization makes use of dependent types in Lean.

Keywords: logic and social choice theory · voting theory · interactive theorem proving · Lean

1 Introduction

There is a long tradition of fruitful interaction between logic and social choice theory. Both Kenneth Arrow [?, p. 154] and Amartya Sen [?, p. 108] have noted the influence of mathematical logic on their thinking about the foundations of social choice theory. Early work using logical methods in social choice theory includes Murakami’s [?] application of results about three-valued logic to the analysis of voting rules, Rubinstein’s [?] proof of the equivalence between multi-profile and single-profile approaches to social choice, and Parikh’s [?] development of a logic of games to study social procedures. There is now a rich literature developing logical systems that can formalize results in social choice theory (see, e.g., [?, ?, ?, ?, ?, ?, ?, ?, ?]).

In recent years, research on logic and social choice theory has centered around the use of computer-aided methods [?]. For example, building on Tang and Lin’s [?] proof of Arrow’s celebrated Impossibility Theorem [?], SAT solvers have

been successfully used to find new impossibility theorems in social choice theory [?, ?, ?, ?, ?], to find minimal examples illustrating how certain voting methods violate key social choice axioms [?], and to study multi-winner voting methods [?] and matching mechanisms [?, ?]. In this paper, we focus on a different line of research on logic and social choice theory that uses interactive theorem provers to verify results in social choice theory. The first applications of interactive theorem provers in social choice used Isabelle/HOL [?] and Mizar [?] to formalize different proofs of Arrow’s Impossibility Theorem [?]. More recently, [?] and [?] used Isabelle to verify theorems from [?] and [?], respectively. These projects demonstrate, as Nipkow [?] notes, that “social choice theory turns out to be perfectly suitable for mechanical theorem proving.” In this paper, we provide further evidence of this by developing a framework for formalizing voting theory using an interactive theorem prover. One obvious benefit of such a project is the verification of the correctness of mathematical claims in voting theory. Several published claims, including Arrow’s [?] original statement of his impossibility theorem (for more than 3 candidates), Baigent’s [?] variation involving “weak IIA” (in the case of 3 candidates), and Routley’s [?] claimed generalization of Arrow’s theorem to infinite populations, were disproved by counterexamples (see [?], [?], and [?]); and who knows what other mistakes might have slipped by referees. Second, formalization allows us to carefully track which assumptions—e.g., about voter preferences, cardinalities, choice of primitive concepts, etc.—are needed for which results, leading to generalizations and perhaps even new avenues for research. Third, formalization may eventually facilitate automated search of the corpus of proved results for use by researchers in proving new results.

For our formalization project we chose to use the Lean theorem prover [?], a framework that supports both interactive and automated theorem proving. Lean’s kernel is based on dependent type theory and implements a version of the calculus of inductive constructions [?] and Martin-Löf type theory [?]. There is an extensive and actively maintained library of mathematical results formalized in Lean (see https://leanprover-community.github.io/mathlib_docs/). In addition, Lean is the system chosen for the Formal Abstracts project initiated by Thomas Hales (<https://formalabstracts.github.io>).

Our aim was to use Lean to verify results about properties satisfied by voting methods (e.g., Condorcet consistency, independence of clones, etc.). In order to formalize voting theoretic axioms concerning adding or removing candidates and voters, we work in a variable-election setting whose formalization makes use of dependent types, as explained in Section 2. In Section 3, we illustrate the usefulness of our framework by verifying almost all the results from [?] about a recently proposed voting method, Split Cycle (defined in Example 3 below). We conclude in Section 5 with directions for further work.

2 Framework

In this section, we define the basic objects of voting theory: profiles, social choice correspondences, etc. We first give standard set-theoretic definitions and then their type-theoretic counterparts in Lean syntax.

2.1 Profiles

For our set-theoretic definitions, we fix infinite sets \mathcal{V} and \mathcal{X} of voters and candidates, respectively. Given $X \subseteq \mathcal{X}$, let $\mathcal{B}(X)$ be the set of all binary relations on X . Instead of thinking of a binary relation as a set of ordered pairs, it is more convenient to think of a binary on X as a function $S : X \times X \rightarrow \{0, 1\}$. In fact, to better match our Lean formalization, we “curry” all functions with multiple arguments, transforming them into functions with single arguments that output functions. Thus, we regard a binary relation on X as a function $S : X \rightarrow \{0, 1\}^X$, where $\{0, 1\}^X$ is the set of functions from X to $\{0, 1\}$. For any $x \in X$, $S(x) : X \rightarrow \{0, 1\}$, and $S(x)(y) = 1$ means that the binary relation S holds of (x, y) . In what follows, we write ‘ xSy ’ instead of $S(x)(y) = 1$.

Definition 1. For $V \subseteq \mathcal{V}$ and $X \subseteq \mathcal{X}$, a (V, X) -profile is a map $\mathbf{Q} : V \rightarrow \mathcal{B}(X)$. We write ‘ \mathbf{Q}_i ’ for the relation $\mathbf{Q}(i)$. Given a (V, X) -profile \mathbf{Q} , let $V(\mathbf{Q})$ be V and $X(\mathbf{Q})$ be X . We then define a function **Prof** that assigns to each pair (V, X) of $V \subseteq \mathcal{V}$ and $X \subseteq \mathcal{X}$ the set $\text{Prof}(V, X)$ of all (V, X) -profiles. Finally, define $\text{PROF} = \bigcup_{V \subseteq \mathcal{V}, X \subseteq \mathcal{X}} \text{Prof}(V, X)$.

Depending on the application, one can interpret $x\mathbf{Q}_iy$ to mean either (i) that voter i strictly prefers x to y or (ii) that voter i strictly prefers x to y or is indifferent between x and y . Under interpretation (i), we use ‘**P**’ for a profile; under interpretation (ii), we use ‘**R**’ for a profile.⁴ A profile \mathbf{Q} is said to be *asymmetric* (*transitive*, etc.) if for every $i \in V$, \mathbf{Q}_i is asymmetric (*transitive*, etc.). Of course, asymmetric profiles only make sense under interpretation (i), whereas under interpretation (ii), profiles should be reflexive.

To translate Definition 1 into Lean, we first think of V and X as types, rather than sets, and then represent the function **Prof** from Definition 1 as follows:⁵

```
def Prof : Type → Type → Type :=
  λ (V X : Type), V → X → X → Prop
```

Here **Prop** is a special type that here plays the role of $\{0, 1\}$ in the formalization of binary relations mentioned above. The definition states that **Prof** is a function that given two types, V and X , outputs the type $V \rightarrow X \rightarrow X \rightarrow \text{Prop}$. Because

⁴ Approach (ii) is more general, since it allows one to distinguish between voter i being *indifferent* between x and y , defined as $x\mathbf{R}_iy$ and $y\mathbf{R}_ix$, vs. x and y being *noncomparable* for i , defined as *neither* $x\mathbf{R}_iy$ *nor* $y\mathbf{R}_ix$. When the distinction between voter indifference and noncomparability is not needed, approach (i) can be simpler.

⁵ When writing type expressions, arrows associate to the right, so, e.g., the expression ‘ $V \rightarrow X \rightarrow X \rightarrow \text{Prop}$ ’ stands for $V \rightarrow (X \rightarrow (X \rightarrow \text{Prop}))$.

$X \rightarrow X \rightarrow \text{Prop}$ is the type of binary relations on X , an inhabitant of the type $V \rightarrow X \rightarrow X \rightarrow \text{Prop}$ can be viewed as a (V, X) -profile. Thus, we may think of $\text{Prof } V \ X$ as the type of (V, X) -profiles.

One of the most important kinds of information to read off from a profile is whether one candidate is majority preferred to another.

Definition 2. Given a profile \mathbf{P} and $x, y \in X(\mathbf{P})$, we say that x is *majority preferred to y in \mathbf{P}* if more voters rank x above y than rank y above x .

In Lean, we formalize Definition 2 as follows:

```
def majority_preferred {V X: Type}: Prof V X → X → X → Prop
:= λ P x y,
  cardinal.mk {v : V // P v x y} > cardinal.mk {v : V // P v y x}
```

Here ‘ $\{V \ X: \text{Type}\}$ ’ indicates that V and X are implicit arguments to `majority_preferred` of type `Type`. The `majority_preferred` function takes in explicit arguments of a (V, X) -profile and two candidates and returns the proposition stating that the cardinality of the set of voters who prefer x to y is greater than the cardinality of the set of voters who prefer y to x .

We often are concerned not only with whether one candidate is majority preferred to another but also, if so, what is the margin of majority preference.

Definition 3. Given a profile \mathbf{P} and $x_1, x_2 \in X(\mathbf{P})$, the *margin of x_1 over x_2 in \mathbf{P}* , denoted $\text{Margin}_{\mathbf{P}}(x_1, x_2)$, is $|\{i \in V(\mathbf{P}) \mid x_1 \mathbf{P}_i x_2\}| - |\{i \in V(\mathbf{P}) \mid x_2 \mathbf{P}_i x_1\}|$.

In Lean, Definition 3 becomes:

```
def margin {V X: Type} [fintype V] : Prof V X → X → X → ℤ :=
λ P x1 x2, ↑(finset.univ.filter (λ v, P v x1 x2)).card -
↑(finset.univ.filter (λ v, P v x2 x1)).card
```

Here ‘ $[fintype V]$ ’ can roughly be understood as an implicit assumption that V is finite,⁶ which we make so that we can perform the subtraction in the definition of margin. The `margin` function takes in explicit arguments of a (V, X) -profile and two candidates and returns the margin of the first over the second; in particular, ‘`finset.univ.filter (λ v, P v x1 x2)`’ is syntax for constructing $\{v \in V(\mathbf{P}) \mid x_1 \mathbf{P}_v x_2\}$, `.card` takes the cardinality of the set (a natural number), and `↑` shifts the type from natural number to integer.

As usual, we can regard the $\text{Margin}_{\mathbf{P}}$ function as an $|X(\mathbf{P})| \times |X(\mathbf{P})|$ matrix. Since $\text{Margin}_{\mathbf{P}}(x, y) = -\text{Margin}_{\mathbf{P}}(y, x)$, the matrix is skew-symmetric. Treating an integer-valued square matrix as a function from a set X to functions from X to \mathbb{Z} , the property of skew-symmetry takes in such a function and outputs the proposition stating the skew-symmetry equation holds for all pairs of elements:

⁶ See Section 3.3 of the Lean documentation and the Lean community page on Sets and set-like objects.

```
def skew_symmetric {X: Type} : (X → X → ℤ) → Prop :=
λ M, ∀ x y, M x y = - M y x.
```

Even such a basic fact as that $\text{Margin}_{\mathbf{P}}$ is skew-symmetric needs to be proved in Lean, but the verification is straightforward using Lean’s automation:

```
lemma margin_skew_symmetric {V X: Type} (P : Prof V X)
[finite V] : skew_symmetric (margin P) :=
begin
  intro,
  intro,
  unfold margin,
  simp,
end
```

As the logical form of the skew-symmetric proposition is $\forall x y, \text{Margin } P \ x \ y = - \text{Margin } P \ y \ x$, we prove the claim by two applications of universal introduction using the `intro` tactic. For arbitrary x and y in X , we must prove $\text{Margin } P \ x \ y = - \text{Margin } P \ y \ x$. This is done by first unfolding the definition of `margin`, so that our goal is now to prove

$$\begin{aligned} & \uparrow(\text{finset.univ.filter } (\lambda v, P \ v \ x \ y)).\text{card} - \\ & \uparrow(\text{finset.univ.filter } (\lambda v, P \ v \ y \ x)).\text{card} \\ & = -(\uparrow(\text{finset.univ.filter } (\lambda v, P \ v \ y \ x)).\text{card} - \\ & \uparrow(\text{finset.univ.filter } (\lambda v, P \ v \ x \ y)).\text{card}) \end{aligned}$$

Fortunately, Lean can prove this equation automatically using the `simp` tactic.

Turning to properties of profiles, one of the most important to consider is whether a profile has a so-called Condorcet winner or even a majority winner.

Definition 4. Given a profile \mathbf{P} with $V(\mathbf{P})$ finite and $x \in X(\mathbf{P})$, we say that x is a *Condorcet winner in \mathbf{P}* if for all $y \in X(\mathbf{P})$ with $y \neq x$, we have $\text{Margin}_{\mathbf{P}}(x, y) > 0$. We say that x is a *majority winner in \mathbf{P}* if the number of voters who rank x (and only x) in first place is greater than the number of voter who do not rank x in first place.

In Lean, Definition 4 becomes:

```
def condorcet_winner (P : Profile V X) (x : X) :
Prop := ∀ y ≠ x, majority_preferred P x y

def majority_winner (P : Profile V X) (x : X) :
Prop := cardinal.mk {v : V // ∀ y ≠ x, P v x y} > cardinal.mk
{v : V // ∃ y ≠ x, P v y x}
```

Here the ‘//’ notation indicates we are identifying the subtype of voters with a certain property.

As an example of a slightly more involved proof than the one above showing that the margin matrix is skew-symmetric, we present a proof in Lean that a majority winner is also a Condorcet winner. For this we use several basic theorems provided by Mathlib, including one formalizing the fact that a subset of a set has cardinality less than or equal to that of its superset:⁷

```
theorem cardinal.mk_subtype_mono {α : Type u} {φ ψ : α → Prop}
(h : ∀ x, φ x → ψ x) :
cardinal.mk {x // φ x} ≤ cardinal.mk {x // ψ x}
```

We explain the following Lean proof that a majority winner is a Condorcet winner in detail below:

```
lemma majority_winner_is_condorcet (P : Profile V X) [fintype V]
(x : X) : majority_winner P x → condorcet_winner P x :=
begin
1.  intros majority z z_ne_x,
2.  have imp1 : ∀ v, (∀ y ≠ x, P v x y) → P v x z,
3.    {intros v ranks_x_first,
4.     exact ranks_x_first z z_ne_x},
5.  refine lt_of_lt_of_le _ (cardinal.mk_subtype_mono imp1),
6.  have imp2 : ∀ v, P v z x → (∃ y ≠ x, P v y x),
7.    {intros v ranks_z_above_x,
8.     use [z, z_ne_x],
9.     exact ranks_z_above_x},
10. apply lt_of_le_of_lt (cardinal.mk_subtype_mono imp2),
11. exact majority,
end
```

Since the logical form of what we want to prove, `majority_winner P x → condorcet_winner P x`, is an implication, we use `intros` on line 1 to introduce a name `majority` for a proof of `majority_winner P x`. Then since the consequent, `condorcet_winner P x`, is a universal statement, $\forall y \neq x$, `majority_preferred P x y`, we introduce a name `z` for an arbitrary candidate and a name `z_ne_x` for a proof of $z \neq x$. Our goal is now to prove `majority_preferred P x z`.

The first key move (lines 4-6) is to prove that everyone who ranks `x` first ranks `x` above `z`. On line 5, we introduce a name `v` for an arbitrary voter and a name `ranks_x_first` for a proof of $\forall y \neq x, P v x y$. Such a proof is a *function* that when given a `w` in the domain of \forall and a proof of $w \neq y$ outputs a proof of $P v x w$. Thus, when given `z` and our proof `z_ne_x`, `ranks_x_first` outputs a proof of $P v x z$, which is exactly what we want, so we can apply Lean's `exact` tactic. Since `imp1` is of the form $(\forall v, \varphi v \rightarrow \psi v)$, we can apply the Mathlib theorem `cardinal.mk_subtype_mono` to obtain a proof `cardinal.mk_subtype_mono imp1` that the number of voters who rank `x` first is less than or equal to the number of voters who rank `x` above `z`.

⁷ We have changed variable names and replaced `#` with `'cardinal.mk'`.

Next, on line 7, we use a Mathlib theorem, `lt_of_lt_of_le`, which states that $n < m \rightarrow m \leq k \rightarrow n < k$ (recall that implication associates to the right). Take n to be the number of voters who rank z above x , m to be the number who rank x first, and k to be the number who rank x above z . Thus, our goal is to prove $n < k$, and above we proved $m \leq k$. Now $m \leq k$ is not the antecedent of $n < m \rightarrow m \leq k \rightarrow n < k$, but Lean’s `refine` tactic allows us to insert a placeholder `_` for the antecedent, so our goal then becomes proving $n < m$.

To prove $n < m$, the key move (lines 9-11) is to prove that everyone who ranks z over x does not rank x first. Then we can apply `cardinal.mk_subtype_mono` to obtain a proof `cardinal.mk_subtype_mono imp2` that the number n of voters who rank z above x is less than or equal to the number—call it m' —of voters who do not rank x first. Thus, we have a proof of $n \leq m'$, so we can apply the implication $n \leq m' \rightarrow m' < m \rightarrow n < m$ provided by the Mathlib theorem `lt_of_le_of_lt` to obtain a proof of $m' < m \rightarrow n < m$. Then since `majority` is exactly a proof of the antecedent of $m' < m \rightarrow n < m$, we obtain a proof of our goal $n < m$, so we are done.

2.2 Functions on profiles

Next we define two kinds of functions that take profiles as inputs. The first, called a *social choice correspondence* (SCC), assign to a given profile a set of candidates, who are considered tied for winning the election. It is common to consider “domain restrictions” on the set of profiles for which the SCC is defined. Thus, one may define an SCC as a function on some set \mathcal{D} of profiles such that for all $\mathbf{Q} \in \mathcal{D}$, we have $\emptyset \neq F(\mathbf{Q}) \subseteq X(\mathbf{Q})$. However, for our purposes, it is more convenient to use the following equivalent approach.

Definition 5. For $V \subseteq \mathcal{V}$ and $X \subseteq \mathcal{X}$, a *social choice correspondence* for (V, X) , or (V, X) -SCC, is a function $F : \text{Prof}(V, X) \rightarrow \wp(X)$. We abuse terminology and call the set $\{\mathbf{Q} \in \text{Prof}(V, X) \mid F(\mathbf{Q}) \neq \emptyset\}$ the *domain* of the (V, X) -SCC.

Let `SCC` be a function that assigns to each pair (V, X) of $V \subseteq \mathcal{V}$ and $X \subseteq \mathcal{X}$ the set of all (V, X) -SCCs.

We represent the function `SCC` in Lean as follows, where `set X` is the type of subsets of X :⁸

```
def SCC := λ (V X : Type), Prof V X → set X
```

The definition states that `SCC` is a function that given two types, V and X , outputs the type `Prof V X → set X`, which is the type of (V, X) -SCCs.

Example 1. For (V, X) , consider the Condorcet SCC for (V, X) defined as follows:

$$\text{Cond}_{(V, X)}(\mathbf{P}) = \begin{cases} \{x\} & \text{if } x \text{ is a Condorcet winner in } \mathbf{P} \\ X(\mathbf{P}) & \text{otherwise} \end{cases}.$$

We represent this (V, X) -SCC in Lean as follows:

⁸ When writing type expressions, function application binds more strongly than arrow, so ‘`Prof V X → set X`’ stands for `(Prof V X) → set X`.

```

def condorcet_SCC {V X: Type} : SCC V X := λ P,
if ∃ x, condorcet_winner P x then {x : X | condorcet_winner P x}
else set.univ

```

The definition states that given a (V, X) -profile \mathbf{P} , if there is a Condorcet winner, then output the set of all Condorcet winners (which will be a singleton), and otherwise output all candidates in X .

Most voting methods (e.g., Plurality, Borda, Instant Runoff) are defined not only for a fixed set of voters and candidates but for any set of voters and candidates, which motivates the following definition.

Definition 6. A *variable-election social choice correspondence* (VSCC) is a function F that assigns to each pair (V, X) of a $V \subseteq \mathcal{V}$ and $X \subseteq \mathcal{X}$ a (V, X) -SCC. We abuse terminology and call the set $\{\mathbf{Q} \in \text{PROF} \mid F(V(\mathbf{Q}), X(\mathbf{Q}))(\mathbf{Q}) \neq \emptyset\}$ the *domain* of the VSCC.

An equivalent but perhaps more intuitive approach would define a VSCC to be a function on PROF (rather than $\wp(\mathcal{V}) \times \wp(\mathcal{X})$) such that for each $\mathbf{Q} \in \text{PROF}$, we have $F(\mathbf{Q}) \subseteq X(\mathbf{Q})$; ⁹ abusing terminology, we could then call the set $\{\mathbf{Q} \in \text{Prop} \mid F(\mathbf{Q}) \neq \emptyset\}$ the *domain* of the VSCC. However, we have presented Definition 6 above because it nicely connects with our formalization in Lean.

In Lean, we define the type of VSCCs as a *dependent function type*:

```

def VSCC : Type 1 := Π (V X : Type), SCC V X.

```

The definition states that a VSCC is a function that for any types V and X returns a function of the type $\text{SCC } V \ X$, i.e., a (V, X) -SCC.

Example 2. We define the Condorcet VSCC as follows, taking advantage of our definition for any V and X of the Condorcet (V, X) -SCC in Example 1:

```

def condorcet_VSCC : VSCC := λ V X, condorcet_SCC

```

The second type of function we consider assigns to a given profile a binary relation on the set of candidates in the profile.

Definition 7. For $V \subseteq \mathcal{V}$ and $X \subseteq \mathcal{X}$, a *collective choice rule for (V, X)* , or (V, X) -CCR, is a function $f : \text{Prof}(V, X) \rightarrow \mathcal{B}(X)$. Let CCR be a function that assigns to each pair (V, X) of $V \subseteq \mathcal{V}$ and $X \subseteq \mathcal{X}$ the set of all (V, X) -CCRs.

Depending on the application, one can interpret the binary relation $f(\mathbf{Q})$ in one of two ways: $xf(\mathbf{Q})y$ can mean (a) x defeats y socially or (b) x defeats or is tied with y socially.¹⁰ Once again, there is also the issue of “domain restrictions.”

⁹ This is the definition of a *voting method* used in [1] with the additional stipulations that $F(\mathbf{Q}) \neq \emptyset$ and that $V(\mathbf{Q})$ and $X(\mathbf{Q})$ are nonempty and finite.

¹⁰ As in Footnote 4, approach (b) is more general, since it allows one to distinguish between “social indifference” and “social noncomparability” (see, e.g., ...). When notions of social indifference and noncomparability are not needed, approach (a) can be simpler.

Under approach (a), we can mark that the CCR is “undefined” on a profile \mathbf{Q} by setting $f(\mathbf{Q}) = X(\mathbf{Q}) \times X(\mathbf{Q})$. Then we can abuse terminology and call $\{\mathbf{Q} \in \text{Prof}(V, X) \mid f(\mathbf{Q}) \neq X(\mathbf{Q}) \times X(\mathbf{Q})\}$ the domain of f . Under approach (b), we can mark that the CCR is “undefined” on \mathbf{Q} by setting $f(\mathbf{Q}) = \emptyset$. Then we can abuse terminology and call $\{\mathbf{Q} \in \text{Prof}(V, X) \mid f(\mathbf{Q}) \neq \emptyset\}$ the domain of f .

A CCR f is said to be *asymmetric* (resp. *transitive*, etc.), if for all \mathbf{Q} in the domain of f , $f(\mathbf{Q})$ is asymmetric (transitive, etc.). Of course, asymmetric CCRs only make sense under interpretation (a) above, whereas under interpretation (b), CCRs should be reflexive.

In Lean, our representation of the function CCR is similar to that of SCC:

```
def CCR := λ (V X : Type), Prof V X → X → X → Prop
```

Example 3. As an example of a CCR, we consider the Split Cycle CCR studied in [1]. The output of the Split Cycle CCR is an asymmetric relation understood as a relation of “defeat” between candidates. A candidate a defeats a candidate b in \mathbf{P} just in case the margin of a over b is (i) positive and (ii) is not equal to the weakest margin in some majority cycle containing a and b . To formalize this definition, we first need a definition of a cycle in a binary relation:

```
def cycle {X: Type} := λ (R : X → X → Prop) (c : list X),
  ∃ (e : c ≠ list.nil), list.chain R (c.last e) c
```

Here the function `cycle` takes in a binary relation R and a list c of elements of X and outputs the proposition stating that (i) there is a proof e that c is not the empty list, and (ii) c is a cycle in R . To express (ii), we use the construction `list.chain R a c`, where R is a binary relation, a is an element of X , and c is a list of elements of X , which means that a is R -related to the first element of c and that every element in the list c is related to the next element in c . Thus, if we take a as the last element of c , this implies that c is a cycle. Applying `c.last` to the proof e that c is not the empty list outputs the last element of c .

Now we are ready to define the Split Cycle (V, X) -CCR:

```
def split_cycle_CCR {V X: Type} : CCR V X :=
  λ (P : Profile V X) (x y : X), ∀ [f: fintype V],
    0 < @margin V X f P x y ∧
    ¬ (∃ (c : list X), x ∈ c ∧ y ∈ c ∧
      cycle (λ a b, @margin V X f P x y ≤ @margin V X f P a b) c)
```

Explain...

Once again, we can consider functions that are not restricted to a fixed set of voters and candidates.

Definition 8. A *variable-election collective choice rule* (VCCR) is a function that assigns to each pair (V, X) of a $V \subseteq \mathcal{V}$ and $X \subseteq \mathcal{X}$ a (V, X) -CCR.

An equivalent but perhaps more intuitive definition takes a VCCR to be a function f on Prof (instead of $\mathcal{V} \times \mathcal{X}$) such that for all $\mathbf{Q} \in \text{Prof}$, $f(\mathbf{Q})$ is a binary relation on $X(\mathbf{Q})$.¹¹ However, we have presented Definition 8 above because it

¹¹ This is the definition of a VCCR used in [1] with the additional stipulation that $V(\mathbf{Q})$ and $X(\mathbf{Q})$ are nonempty and finite.

nicely connects with our formalization in Lean, which as before defines the type of VCCRs to be a dependent function type:

```
def VCCR :=  $\Pi$  (V X : Type), CCR V X.
```

The definition states that a VCCR is a function that for any types V and X returns a function of the type $\text{CCR } V \ X$, i.e., a (V, X) -CCR.

Example 4. We define the Split Cycle VCCR as follows, taking advantage of our definition for any V and X of the Split Cycle (V, X) -SCC in Example 3:

```
def split_cycle_VCCR : VCCR :=  $\lambda$  V X, split_cycle_CCR
```

3 Theorems

4 Discussion

5 Conclusion

6 First Section

6.1 A Subsection Sample

Please note that the first paragraph of a section or subsection is not indented. The first paragraph that follows a table, figure, equation etc. does not need an indent, either.

Subsequent paragraphs, however, are indented.

Sample Heading (Third Level) Only two levels of headings should be numbered. Lower level headings remain unnumbered; they are formatted as run-in headings.

Sample Heading (Fourth Level) The contribution should contain no more than four levels of headings. Table 1 gives a summary of all heading levels.

Table 1: Table captions should be placed above the tables.

Heading level	Example	Font size and style
Title (centered)	Lecture Notes	14 point, bold
1st-level heading	1 Introduction	12 point, bold
2nd-level heading	2.1 Printing Area	10 point, bold
3rd-level heading	Run-in Heading in Bold. Text follows	10 point, bold
4th-level heading	<i>Lowest Level Heading.</i> Text follows	10 point, italic

Displayed equations are centered and set on a separate line.

$$x + y = z \tag{1}$$

Please try to avoid rasterized images for line-art diagrams and schemas. Whenever possible, use vector graphics instead (see Fig. 1).

Theorem 1. *This is a sample theorem. The run-in heading is set in bold, while the following text appears in italics. Definitions, lemmas, propositions, and corollaries are styled the same way.*

Proof. Proofs, examples, and remarks have the initial word in italics, while the following text appears in normal font.

For citations of references, we prefer the use of square brackets and consecutive numbers. Citations using labels or the author/year convention are also acceptable. The following bibliography provides a sample reference list with entries for journal articles [?], an LNCS chapter [?], a book [?], proceedings without editors [?], and a homepage [?]. Multiple citations are grouped [?,?,?], [?,?,?,?].




Fig. 1: A figure caption is always placed below the illustration. Please note that short captions are centered, while long ones are justified by the macro package automatically.